

# K4 Writeup

## 1 Summary

This is a writeup for Kernel Assignment 4 of CS452, where we have added in interrupt-based I/O, including large amounts of low-level register interaction using this new I/O, and recreated the functionality of A0, including train commands and switch commands. In this writeup, we will describe the structure of the kernel, talk about the choices we made for critical system parameters and limitations, and also talk about the kernel's runtime behaviour and performance measurements.

## 2 Kernel Structure

### 2.1 Kernel Class + Augmentations

#### 2.1.1 K1 Contributions

The main component of our kernel is the `Kernel` class:

```
class Kernel {
public:
    Kernel();
    ~Kernel();
    void schedule_next_task();
    void activate();
    void handle();
    void handle_syscall();
    void handle_interrupt(InterruptCode icode);
    void start_timer();

private:
    int p_id_counter = 0;
    int active_task = 0;
    InterruptFrame* active_request = nullptr;

    Task::Scheduler scheduler;
    Descriptor::TaskDescriptor* tasks[Task::USER_TASK_LIMIT] = { nullptr };
    SlabAllocator<Descriptor::TaskDescriptor, int, int, int, void (*)()> task_allocator
        = SlabAllocator<Descriptor::TaskDescriptor, int, int, int, void (*)()>(
        (char*)Task::USER_TASK_START_ADDRESS, Task::USER_TASK_LIMIT);
    Clock::TimeKeeper time_keeper = Clock::TimeKeeper();

    // clock notifier "list", a pointer to the notifier
    int clock_notifier_tid = Task::CLOCK_QUEUE_EMPTY;

    void allocate_new_task(int parent_id, int priority, void (*pc)());
    void handle_send();
    void handle_receive();
    void handle_reply();
    void handle_await_event(int eventId);

    int idle_tid = SystemTask::IDLE_TID;
};
```

This class wraps the kernel functionality and data structures, so it's a useful reference point for talking about structure and functionality. Most of the class contents are straightforward or self-explanatory, so we will only cover the components that hold nontrivial complexity.

Scheduling: To schedule the order of user tasks, the kernel defers to the `Scheduler`, with `schedule_next_task()` also making use of the class functionality:

```
#define NUM_PRIORITIES 8

class Scheduler
{
public:
    Scheduler();
    int get_next();
    void add_task(int priority, int task_id);

private:
    RingBuffer<int> ready_queue[NUM_PRIORITIES];
};
```

The scheduler itself is very simple, being little more than a wrapper around 3 stack-allocated ring buffers (the ring buffers themselves are a template class with a fixed size of 512 slots; 512 was chosen because we reasoned this would be large enough for any use case). Each buffer stores task IDs and corresponds to a different priority level – high, medium and low – where each queue can be added to on request, and `get_next()` will loop through queues in order to determine which queue should be popped from when called. 8 priorities was determined to be sufficient for our needs (as of K3).

Task Storage: The kernel keeps track of a `SlabAllocator` structure, which keeps track of slab allocations. It is a templated class with variadic arguments, which allows it to construct arbitrary classes in slabs, and it internally uses a ring buffer to keep track of free memory locations. For user task descriptors, slab allocation is set to start at the hardcoded address of `0x10000000`, which we determined to be a safe address for slab allocations (will not collide with important memory addresses).

```
template <typename T, typename... Args>
class SlabAllocator {
public:
    SlabAllocator(char* starting_location, int total_slabs);
    ~SlabAllocator();
    T* get(Args... arguments);
    void del(T* target);
    int get_remaining_size();

private:
    int size;
    int T_size;
    RingBuffer<char*> slabs;
};
```

The slab allocator maintains basic free and delete functions, allowing it to perform similar functionality to that of a heap.

Task descriptors store the task ID, task parent ID, the stack of the task that they represent and some other useful status information.

Activate/Handle: These functions are called in our `kmain` function, which schedules tasks and calls these two functions in an infinite loop. `activate()` initializes user tasks if they have not been initialized and context switches into them, while `handle()` responds to system calls using the kernel-defined `HandlerCodes`.

### 2.1.2 K2 Contributions

The main architecture contributions of K2 are task state and message passing functionality. To support this, task descriptors have been supplemented with some new functions and data structures:

```
class TaskDescriptor {
public:
    enum TaskState { ERROR = 0, ACTIVE = 1, READY = 2, ZOMBIE = 3, SEND_BLOCK = 4, ... };
    TaskDescriptor(int id, int parent_id, int priority, void (*pc)());
    // message related api
    void queue_message(int from, char* msg, int message_length); // queue_up a message
    bool have_message();
    int fill_message(Message msg, int* from, char* msg_container, int msglen);
    int fill_response(int from, char* msg, int msglen);
    Message pop_inbox();

    InterruptFrame* to_active();
    void to_ready(int system_response, Scheduler* scheduler);
    bool kill();
    void to_send_block(char* reply, int replylen);
    void to_receive_block(int* from, char* msg, int msglen);
    void to_reply_block();
    void to_reply_block(char* reply, int replylen);

    bool is_active();
    bool is_ready();
    bool is_zombie();
    bool is_send_block();
    bool is_receive_block();
    bool is_reply_block();

    const int task_id;
    const int parent_id; // id = -1 means no parent

    friend class Kernel;
protected:
    void show_info();
private:
    TaskState state;
    int priority;
    int system_call_result;
    bool initialized;
    void (*pc)();
    MessageReceiver response;
    RingBuffer<Message> inbox;
    char* sp;
    char* kernel_stack[USER_STACK_SIZE];
};
```

These different functions can place tasks into different states, which may or may not remove them from scheduling or flag them as having passed into different states of message passing. Tasks also now maintain a message inbox, which allows them to receive messages even if they aren't actively waiting for them.

As of K2, `USER_STACK_SIZE` is equal to 131,072, equivalent to 128kB of stack space. This was raised up from 16kB (previously 4kB) to make room for an unordered map for the name server.

The message passing functions themselves, then, are responsible for putting tasks into different states, as well as performing memcpy on the messages to transfer them to the right spots. Here's the `handle_send()` function, as an example.

```
void Kernel::handle_send() {
    int rid = active_request->x1;
    if (tasks[rid] == nullptr) {
        // communicating a non existing task
        tasks[active_task]->to_ready(NO_SUCH_TASK, &scheduler);
    } else {
        char* msg = (char*)active_request->x2;
        int msglen = active_request->x3;
        char* reply = (char*)active_request->x4;
        int replylen = active_request->x5;
        if (tasks[rid]->is_receive_block()) {
            tasks[rid]->fill_response(active_task, msg, msglen);
            tasks[rid]->to_ready(msglen, &scheduler);
            tasks[active_task]->to_reply_block(reply, replylen);
        } else {
            // reader is not ready to read we just push it to its inbox
            tasks[rid]->queue_message(active_task, msg, msglen);
            tasks[active_task]->to_send_block(reply, replylen);
        }
    }
}
```

Beyond this, we also create two servers using these message passing functions: the name server and the Rock/Paper/Scissors server. Both servers wait for send requests in an infinite loop, and process decisions only when their receive calls are fulfilled.

- The name server handles name registration and name lookup requests through the `WhoIs()` and `RegisterAs()` system calls. It uses a templated hashmap to keep track of names and task IDs, where names are mandated to be at most 16 characters in length (longer names are truncated). Duplicate names in `RegisterAs()` overwrite the the name server's mapping, while unregistered names in `WhoIs()` cause it to return a pre-specified error code.
- The Rock/Paper/Scissors server allows special clients to sign up for and play in Rock/Paper/Scissors matches against each other by sending special messages to the server. It mediates interactions between clients, keeps track of matches and results and even returns some output showing the results. This will be further discussed in the Kernel Output section.

The templated hashmap used by the name server uses a generic FNV hash function with separate chaining to resolve collisions, and is courtesy of the Embedded Template Library for C++.

### 2.1.3 K3 Contributions

K3 adds interrupt handling, clock functionality and some client and idle tasks, which test these functionalities and track how much time the kernel is spending idle.

**Interrupt Handling:** Besides the assembly-based interrupt handler, managing interrupts also requires slightly different user re-entry due to the volatility of interrupts. Interruption status is tracked via a bool in the task descriptor, and when returning to user tasks, this bool is used to determine if we want syscall re-entry (optimized for ABI rules) or interrupt re-entry (assume nothing).

**Clock Functionality:** To keep track of the clock, clock interrupts, clock ticks, idle time calculations and clock system calls, we introduce the `TimeKeeper` class along with two user tasks: the clock server and the clock notifier. Upon initialization, the kernel enables interrupts, then instantiates an instance of the `TimeKeeper` class, which activates the compare registers and enforces that they occur every 10ms.

- The clock server is responsible for handling the three clock-related system calls: `Time()`, `Delay()` and `DelayUntil()`. Upon receiving a request, the clock server will either reply with an internally-tracked tick integer or add the task ID to an internally held priority queue (of maximum size 64, which is enough to keep track of hopefully enough tasks but not so many as to cause excessive cache misses), which is checked every tick to determine if there is a task that needs to be awoken from delay.
- The clock notifier is awoken whenever a timer interrupt is fired, which is handled using `AwaitEvent`. After being woken up, the clock notifier sends a message to the clock server telling it a tick has occurred, then calls `AwaitEvent` again.
- Here's what the `TimeKeeper` looks like:

```
class TimeKeeper {
public:
    TimeKeeper();
    ~TimeKeeper();

    void start();
    void tick();
    void calculate_and_print_idle_time(int active, int prev, int idle);

private:
    void set_comparator(uint32_t interrupt_time, uint32_t reg_num = 1);
    uint64_t tick_tracker = 0;

    // Time tracking variables
    uint64_t idle_time = 0;
    uint64_t last_ping = 0;
    uint64_t total_time = 1; // start at 1 to avoid division by 0 errors
    uint64_t last_print = 0;
};
```

Here `start()` begins timer tracking, `tick()` advances the internal clock by one tick (10ms) and resets the comparator, and `update_time()` updates idle time calculations (and possibly prints the results to the terminal).

- Client and idle tasks: they work pretty much as specified in the K3 assignment spec. Clients ask the first user task for a given delay and repeat amounts, then print whenever their delays expire. The idle task, which has a lower priority than all clients, simply yields in a loop.

### 2.1.4 K4 Contributions

The K4 contributions can be split into three categories:

- **UART Interfacing:** The Objective of UART interfacing is to avoid busy waiting as much as possible, a.k.a we only read if hardware tell us data is available (as well as we want to read). The work done requires working with very low-level register primitives, talking directly to the UART cable and making use of general-purpose IO (GPIO) and special UART interrupt functionality to allow the kernel to manage I/O in an interrupt-based fashion.

In-order to enable interrupt, we permanently opens IRQ 145 which correspond to GPIO 24, where both UART channel is able to communicate with us through interrupt. This naturally causes a problem: every uart interrupt imaginable comes from the same pipeline, 145, thus, we need to check special register, *IER* to ensure everything is cleared. If we are ever interrupted through IRQ 145, we will continuously read and clear / disable interrupt indicated by *IER* register until both uart0 and uart1's *IER* returns 0x1, which indicate all interrupt is cleared.

As we may have mentioned, sometimes we clear interrupt, sometimes we disable interrupt. define **clearing** as the action disable the interrupt naturally, while **disable** as the action to tell *IER* register to disable certain interrupt. both method have upside and downside. the **clearing** approach ensures that you are interrupted as soon as possible, and avoid missing potential interrupt ticks and fall behind. For example, we relies on *MSR* register and the modem interrupt to read and handle *CTS* interrupt, which calls us every time the level of *CTS* changes. This action need to be handled urgently, since going up (cannot write) and going back down (can write again) may arrive very closely to each other, the risk of turning off the interrupt and potentially miss a tick is just not acceptable. (this can permanently killed the uart1 server until we introduce timeout). However, this approach is not universally viable, since sometimes you have to disable interrupt since there is no way to clear it. For example, the *THR*, which indicate the transmit interrupt cannot be cleared unless you fill the buffer above a certain threshold again. This is either difficult or impossible to do, thus, it is better to shut them down completely during kernel and just let the corresponding listeners transfer the message to kernel.

However, it is pretty good that most interrupt we have to handle either won't have a risk of skipping if we disable, or have a risk of skipping but can be handled by simply clearing the interrupt.

It also worth mentioned that we are trying our RTOS as a hard RTOS, meaning skipping an interrupt is unacceptable behaviour and is considered as system failure. So for, to listen to each interrupt, we only have 1 listener for each, but our code is fast enough such that it totally fine (like, we are at 98% idle)

We chose to make use of a four-server architecture, one for each of the combinations of UART channel 0/1 and input/output, in the hopes that this would lead to optimal overall performance due to servers not needing to block on each other's workflows.

Notably, this particular architecture choice takes advantage of the fact that each of the UART use cases is relatively specialized. UART0 input and output are used exclusively by terminal tasks, which in our system we have confined to one task each, so the server usage can be tightly controlled. UART1 input is also used exclusively for reading sensor data, so specializing the server for that purpose – specifically, the purpose of reading 10 bytes every 100ms (in our system) – makes for a more optimized structure overall. UART1 output, which controls all train command functionality, requires a bit more care, although setting it to its own server gives us more modularization and control regardless.

As for the train and sensor, we decided to split them apart. Train server mostly handles the train command and the track command which are mostly input, while the sensor server is mostly reading output from uart1. this splitting of workflow allow an almost complete decouple of workflow between input and output, while allowing any server that relies on these server to talking to one of them have no influence on the other (except on the wire physically but is outside the scope of K4 for this case).

The Train server either pump corresponding byte right away to the uart1 write server, or create a

task that pump some control byte to the uart1 server after some delay. However, there is no much write protection right now (for example, train is allow to actually write speed command which will be overwrite by reverse command's later acceleration after delay), but they will be enforced in TC1.

The Sensor server have a worker that just keeps pulling from uart1 read server with a 100 ms delay (though we are not sure if we even need it cause MMU and data cache actually makes our code go flying, only concern is the 2400 bauds rate and the fact that the wire is 1 directional). Any task can subscribe to the sensor server, which will be updated as soon as the new reading comes in. The workflow is pretty much (Sensor worker detect update → notify Sensor admin → notify all subscriber → Sensor worker listen for update). This allow the most up to date information for all subscribers, and if any subscriber want continuous update, it should re-subscribe right after getting back as well.

Following is the core of how we handle uart interrupt, though not pretty, it works pretty well. (we will be extracting this feature to a specific subclass that records and handles the interrupt) The text is pretty small, so I recommend reading the actual branch if you are interested in more details.

```
case InterruptCode::UART: {
    /**
     * Note that no matter which interrupt, you receive from the same id, UART_INTERRUPT_ID
     * this is kinda a problem, since a large variety of stuff can be from that type of interrupt,
     * even same type, but uart0 vs uart1
     *
     * in order to differentiate them, we relies on checking the register later, IIR,
     * IIR includes both the information about if there is an interrupt to handle, and if so, what is the interrupt exactly.
     *
     * the general work flow is 1. we receive UART_INTERRUPT_ID interrupt, thus recognize interrupt happened
     * we check IIR to see what type of interrupt happened, we could potentially get a large quantity of interrupt
     * overall, the goal is that if we receive interrupt in the form of UART_INTERRUPT_ID, we keep clearing interrupt
     * until every interrupt associated with uart is cleared
     *
     * Also note that server is in control of which register is flipped, thus also in control of which interrupt is happening.
     */

    int exception_code = (int)(uart_get(0, 0, UART_IIR) & 0x3F);
    do {
        if (exception_code == UART::InterruptType::UART_RX_TIMEOUT && uart_0_receive_tid != Task::UART_RECEIVE_EMPTY) {
            int input_len = uart_get_all(0, 0, tasks[uart_0_receive_tid]→get_event_buffer());
            tasks[uart_0_receive_tid]→to_ready(input_len, &scheduler);
            uart_0_receive_tid = Task::UART_RECEIVE_EMPTY;
            enable_receive_interrupt[0] = false;
            interrupt_control(0);
        } else if (exception_code == UART::InterruptType::UART_TXR_INTERRUPT && uart_0_transmit_tid != Task::UART_TRANSMIT_FULL) {
            tasks[uart_0_transmit_tid]→to_ready(0x0, &scheduler);
            uart_0_transmit_tid = Task::UART_TRANSMIT_FULL;
            enable_transmit_interrupt[0] = false;
            interrupt_control(0);
        } else if (exception_code == UART::InterruptType::UART_CLEAR) {
            break;
        } else {
            printf("UART 0 Too Slow \r\nexception code: %d receive_tid: %d transmit_tid %d\r\n", exception_code, uart_0_receive_tid, uart_0_transmit_tid);
            while (true) {
            }
        }
        exception_code = (int)(uart_get(0, 0, UART_IIR) & 0x3F);
    } while (exception_code != UART::InterruptType::UART_CLEAR);

    exception_code = (int)(uart_get(0, 1, UART_IIR) & 0x3F);
    do {
        // this is a really shitty way to handle this, I think it would probably be better if we something similar to a dedicated class object
        // but we will fix it soon once experienta go through.
        if (exception_code == UART::InterruptType::UART_RX_TIMEOUT && uart_1_receive_timeout_tid != Task::UART_RECEIVE_EMPTY) {
            tasks[uart_1_receive_timeout_tid]→to_ready(0x0, &scheduler);
            uart_1_receive_timeout_tid = Task::UART_RECEIVE_EMPTY;
            enable_receive_interrupt[1] = false;
            interrupt_control(1);
        } else if (exception_code == UART::InterruptType::UART_RX_INTERRUPT && uart_1_receive_tid != Task::UART_RECEIVE_EMPTY) {
            tasks[uart_1_receive_tid]→to_ready(0x0, &scheduler);
            uart_1_receive_tid = Task::UART_RECEIVE_EMPTY;
            enable_receive_interrupt[1] = false;
            interrupt_control(1);
        } else if (exception_code == UART::InterruptType::UART_MODEM_INTERRUPT && uart_1_msr_tid != Task::UART_TRANSMIT_FULL) {
            char state = uart_get(0, 1, UART_MSR);
            if ((state & 0x1) == 0x1) {
                tasks[uart_1_msr_tid]→to_ready(0x0, &scheduler);
                uart_1_msr_tid = Task::UART_TRANSMIT_FULL;
            }
        } else if (exception_code == UART::InterruptType::UART_TXR_INTERRUPT && uart_1_transmit_tid != Task::UART_TRANSMIT_FULL) {
            tasks[uart_1_transmit_tid]→to_ready(0x0, &scheduler);
            uart_1_transmit_tid = Task::UART_TRANSMIT_FULL;
            enable_transmit_interrupt[1] = false;
            interrupt_control(1);
        } else if (exception_code == UART::InterruptType::UART_CLEAR) {
            break;
        } else {
            printf("UART 1 Too Slow \r\nexception code: %d receive_tid: %d transmit_tid %d msr_tid %d\r\n", exception_code, uart_1_receive_tid, uart_1_transmit_tid, uart_1_msr_tid);
            while (true) {
            }
        }
        exception_code = (int)(uart_get(0, 1, UART_IIR) & 0x3F);
    } while (exception_code != UART::InterruptType::UART_CLEAR);
    UART::clear_uart_interrupt();
    break;
}
```

- Terminal management: The ultimate end-user functionality requirement of K4 is to have recreated the functionality of A0, which necessarily requires a considerable amount of terminal management. Specifically, we need to gather and display information about system time, idle time, sensors and switches, and also allow users to control train speeds, train directions and switch layouts through a command line interface. To recreate this, a dedicated user input "courier" is used, which is the only task that reads from UART0, and a dedicated terminal server is used, which is the only task that is allowed to write to UART0. All UART0 writing requests must go through the terminal server, which carefully controls and formats output.
- User task architecture: On a more general level, the layout of the user tasks as a whole has been carefully designed and planned out to weave a complicated net of tasks, making heavy use of the ideas of servers, couriers and workers to offload work to other tasks if necessary and make sure nothing is needlessly blocked. Here's a rough sketch of what the current system architecture looks like:

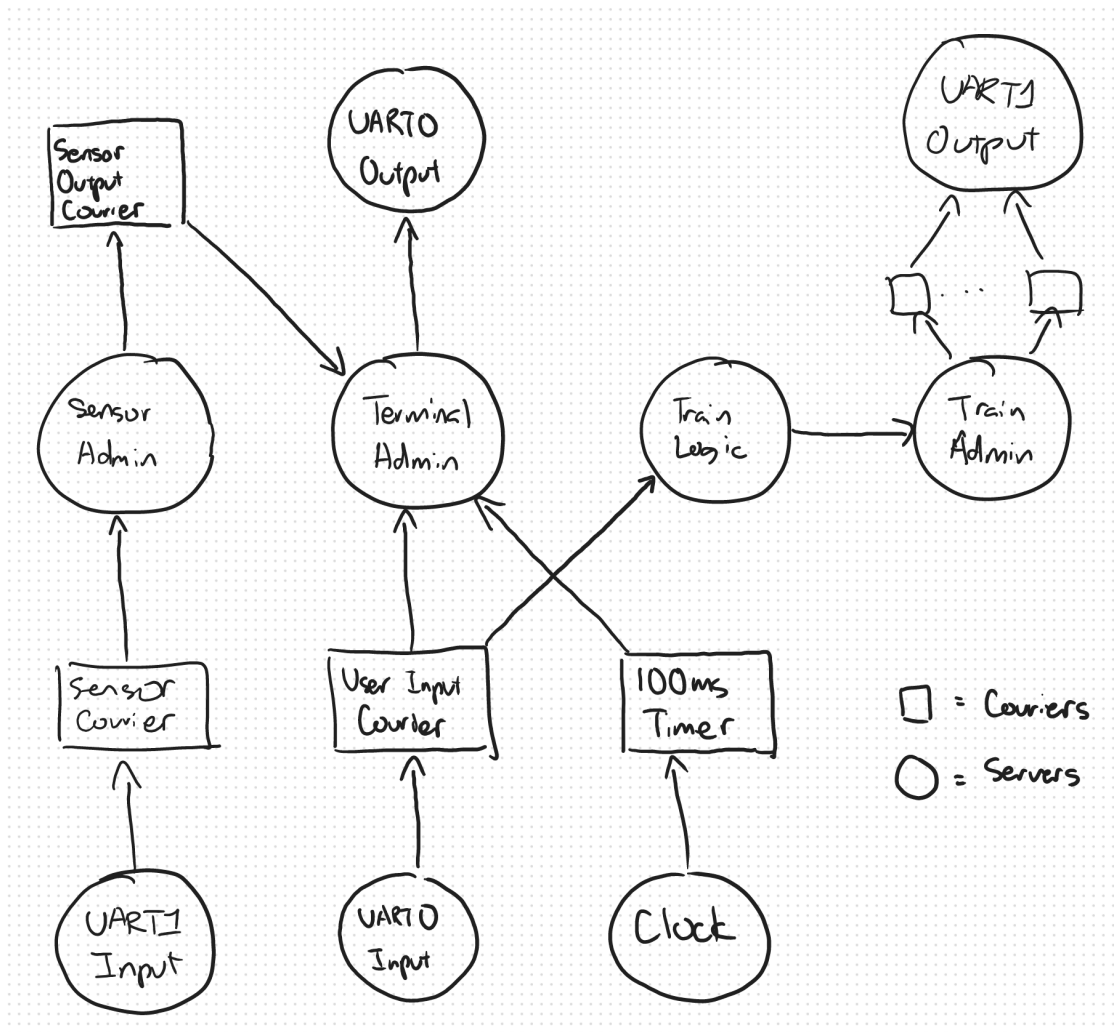


Figure 1: A highlighted subset of our K4 architecture additions. Not shown: the idle task/idle timer task.

In general, couriers are used when information needs to be passed between tasks without blocking receiving. The train admin server in particular makes use of a courier pool to service different train commands without becoming blocked itself, especially as some train commands may be quite slow to process due to the slow speed of train command listening and/or UART1 transmission.



## 2.2 Context Switching

Most of our context switching procedure is relatively standard/inflexible, so I'll just jot down some key features here:

- For the sake of optimization, we maximally assume ABI rules hold when doing system call context switching. Indeed, we can even disregard the value of the program counter, since calling the system call means that the link register now holds the correct value to return to. On the other hand, the interrupt handler saves and loads all registers when entering/exiting the kernel, and also does a bit of register juggling to keep track of the program counter and stack pointer.
- When the context switch returns to the kernel, it passes the kernel a pointer to an `InterruptFrame` struct, which is essentially a format of the 31 main non-zero registers, allowing easy access, along with 3 extra 8-byte chunks of information: the program state of the user, read from `SPSR_EL1`; the program counter of the user, read from `ELR_EL1`; and a data chunk, which indicates that the user task was interrupted rather than performing a system call.
- The first time a user task is switched into is different from subsequent switchings, as registers do not need to be loaded the first time.
- The `x0` register is used liberally to pass around arguments and return values.

## 2.3 RPS Server/Clients (K2 only)

The RPS Server, like the name server, handles sends/receives from RPS clients, coordinates match logic and keeps track of matches (storing matches in a large, linearly-searched array).

Clients loop through the following logic loop:

1. Ask for the task ID of the RPS server, and send it a signup request.
2. Query the current system time to make a "random" play decision, and send it to the server.
3. Repeat the previous step  $3 + \text{randint}(0, 2)$  times, then send a quit message.
4. After sending a quit message, or after receiving a quit message from the server, either destroy self (75% chance) or rejoin the match pool (25% chance).

## 2.4 Idle/Client Tasks (K3 only)

In K3, we are responsible for setting up the following user task scheme:

1. The first user task creates the name server, clock server, clock notifier, then sets up 4 client tasks with priorities 3, 4, 5 and 6 and an idle task with priority 7, then waits for the 4 client tasks to send requests. Upon receiving, the user task sends out two integers: a delay amount and a repetition count, encoded in a string.
2. The client tasks, upon creation, ask the first user task for a delay amount  $D$  and a repetition count  $C$ . They then send  $C$  requests to the clock server to delay by  $D$  ticks, printing every time they finish a delay (or rather, ask the kernel to print on their behalf, as user-level printing was determined to be unstable due to interruption), and exit on completion.
3. The idle task simply yields repeatedly in an infinite loop.

### 3 Memory Management Unit

Though MMU itself is not spectacular, it is the required hardware configuration if we want to enable some form of data caching. Thus, the minimum requirement is a translation table that flat maps address directly to each entry on the table.

we will not talk about many details related to MMU itself, but more about what design decision that we thought works the best. If I go into detail why certain parameter are set a certain way, this will take about 10 pages to complete.

1. We used a level 1 table and four level 2 table for flat-mapping. Once again, if the goal is to maximize performance through stuff like TLB we might have to go in deeper, but for our purpose this is enough
2. since the only interesting memory address are the first 4 gb (first 2gb is real, while device memory is located somewhere in the 4th gb). the goal is to map them according to their usage
  - (a) Our code is located at 0x80000, the default start location of pi-4, and for convenience reason, we will map the entire first 2mb block from 0x0 to 0x100000 as execution memory, giving el1 the permission to write/read/execute and el0 the permission to only execute. Linker will ensure this area only contains code which we execute. note that our current executable is only about 46 kb when it comes to .text section, thus this area should be more than enough to contain all the code
  - (b) Future byte all the way until 2gb mark is simply marked as memory that can be read/write for both exception level, in the linker I assured any .data or related content would be moved out of this area. This include global constants and user stacks. We also choose both inner and outer write back Non-transient, this is because we seems to easily max out the L1 cache on processor 0, and most of our operation could be contained within the 2mb storage provided by L2 cache. this means very little use of actual memory is used, so we don't have to worry about writing back to memory (also there is no multi core so there is no cache coherency).
  - (c) Some part of the memory are device specific memory. They are not real memory addresses, but the physical location will magically convert processor access into the right register. Thus, this area is similar to marking as read/write for both exception level, but we need to avoid caching. If a write to register is cached, it essentially have 0 effect.
3. it also worth noting that we intentionally turned off TCBR1\_EL1 which is suppose to point toward the upper address beyond the 4GB, but since we will never access that region in correct execution sequence, it nice to turn it off. in case of a bug, we will at least get some format of exception.
4. we didn't have a level 0 table because there is no need to support one, the level 1 table already support up to 512 gb of memory mapping, and I am only using 4gb. (unfortunately level 2 table only support 1gb so I couldn't go further, however, there are some setting you can tune such that table will only use 4 entries)
5. funny enough, the 3rd table (for the 3rd gb) is completely irrelevant to our stuff, there is no real memory there and there is no device memory that relies on it, so we are free to actually just, not even initialize it and leave it blank, but I kinda just filled it in anyway

that is pretty much most of the design, later on we simply setup the related parameter to true and flip then switch on STCLR\_EL1 and everything would start running.

Lastly, we stashed the table starting from 0x40000, and point TCBR0\_EL1 toward it then we have a properly working MMU, though just flat mapping. With MMU, we have dcache, with dcache, our SRR time went from 31 micro seconds to almost < 1 microseconds.

## 4 Kernel Output (K4)

Here is a sample of what typical output from our kernel looks like:




Figure 2: Screen capture of sample K4 output.

Here's what each of the different lines mean. Note that some of these lines do not appear in the figure above because we have now separated our kernel into two phases: a bootup phase, which prints some helpful logging, and a command line phase, which is entered after typing any key after bootup and which is shown in the figure above. Only in the command line phase can commands actually be entered.

- **init kernel:** The `kmain` function has been called.
- **finished kernel init, started scheduling user tasks:** The Kernel class has been created.
- **mmu setup complete:** The MMU has finished being setup, with memory mapping and table creation.
- **MMM:SS:T:** The timer, in the top left, showing current system time in tenths of a second.
- **Idle time: <Ti>:** The number of microseconds that have passed while inside the idle task (roughly).
- **Total time: <Tt>:** The total number of microseconds that have passed since kernel initialization.
- **Percentage: <P>:** The percentage of idle time over total time.
- **SENSOR DATA:** A list of recently triggered sensors will appear underneath this line in red.
- **SWITCHES:** The switch table, showing the status of all switches. There is also a line of '???' characters embedded in the table – this is a printing bug that, for some reason, only appears during initialization. We're not sure of its source, but it seems to be related to previous UART0 output, as though switch initialization causes some kind of buffer flush (if we remove switch initialization, this line of '???' characters disappears).
- **ABYSS> <command>:** The status of the terminal after a command has been entered. Note the three different main commands, as well as the quit command, which also sets the speed of all trains to 0.

A couple other known bugs: for some reason, sometimes if input is entered in too fast, random 'a' characters will appear on the screen. Also, resetting the Märklin box can cause unhandled UART1 exceptions, and the reverse command may be inconsistent on train 2 (possibly because it is the fastest train by a significant margin).

## 5 Acknowledgements

Special thanks to:

- Mike Krinkin for his excellent ARMv8 tutorial, available [here](#). Without it, writing a stable context switch would likely have been much more difficult.
- John Wellbelove and other contributors for their work on the [Embedded Template Library \(ETL\)](#) which we have shamelessly lifted directly for use in our kernel thanks to the wealth of data structures it provides.
- Marco Paland for [his work](#) on reimplementing `printf`, a function that is incredibly useful for debugging purposes, and requiring only a simple `putchar()` implementation to function properly.
- Rock Zhang (@codingbelief on GitHub) for his [comprehensive ARMv8 tutorial](#), which was invaluable for setting up the MMU.
- Darwin Chen for being the brave soul who first plundered into the depths of MMU configuration territory, and who provided hope in dark times when things weren't working properly.